

KAKASHI

守

LOCAL
MINIMAL
VERIFIED

Manager briefing • capabilities, government value, and live
proof • v1.2.0

MANAGER BRIEFING • LIVE DEMO

The risk has changed

- AI agents read files, query systems, and call tools before a human reviews every payload.
- A single attachment can contain government identifiers, contact data, finance records, or production credentials.
- Once raw data reaches an external model or service, the disclosure cannot be undone.
- “Remember to redact it” is guidance. Government deployments need an enforceable, testable control.

THE CONTROL POINT
Kakashi stands between
local data and the
destination.

Detect → Minimise →
Decide → Verify

Three layers, one local trust boundary

- ENGINE — detect and transform sensitive values in files while preserving the original format.
- REACH — apply the same protection to folders, repositories, shared drives, and database query results.
- GUARDIAN — reason about the agent, declared task, destination, and policy before releasing anything.
- Each layer is useful independently; together they form a privacy gate for agentic workflows.

1 ENGINE
2 REACH
3 GUARDIAN

All processing stays on
the machine.

Coverage designed for real government work

- 35 active patterns across credentials, government IDs, financial data, contacts, names, and technical identifiers.
- 50+ formats: PDF, Word, PowerPoint, Excel, CSV, structured data, configuration, and source code.
- Typed tokens preserve relationships; redact removes values; fake mode preserves realistic shape.
- English and Arabic CLI/report experiences support bilingual operational teams.

35
DETECTION PATTERNS

50+
FILE FORMATS

EN • العربية

Case 1 — Protect a file before an AI receives it

- Scan is count-only by default: an agent learns the risk without seeing matched values.
- Audit is explicit and human-only because it intentionally reveals the original-to-token mapping.
- Mask writes a safe sibling rather than overwriting the source.
- The demo proves typed, redact, fake, whitelist, stdin, and re-scan behaviour.

LIVE COMMAND
./demos/01-file-
protection.sh

PROOF
Source values removed
Typed/redact outputs re-
scan clean

Case 2 — Discover risk across a data estate

- Recursive scanning respects repository ignore policy and reports files intentionally skipped.
- JSON, Markdown, and HTML reports contain finding metadata without raw matched values by default.
- PDPL-oriented article mapping gives compliance teams reviewable context—not a legal certification.
- Arabic HTML uses the same local evidence and renders without a cloud translation step.

LIVE COMMAND
./demos/02-folder-compliance.sh

OUTPUT
English HTML
Arabic HTML
JSON + Markdown

Case 3 — Mask database results client-side

- Adapters support PostgreSQL, MySQL, MongoDB, Snowflake, Databricks, and SQLite.
- Rows are transformed locally before a safe JSONL, JSON, or CSV copy is written.
- Stable tokens preserve distinct counts and relationships across rows for useful analytics.
- Limits reduce accidental bulk access; source databases are not modified by the export flow.

LIVE COMMAND
./demos/03-database-
protection.sh

DEMO ADAPTER
Offline mock dataset
No database required

Case 4 — Guardian makes the release decision

- Observe → understand task → assess risk → plan → policy check → act → verify → replan.
- Decisions are explicit: ALLOW, ALLOW_WITH_TRANSFORMATION, REQUIRE_APPROVAL, or BLOCK.
- A verified analytics request receives only the transformed release path.
- A credential-bearing request fails closed and releases no artifact without human approval.

LIVE COMMAND
./demos/04-guardian-
decision.sh

KEY PROPERTY
The agent cannot self-
approve.

Case 5 — A standing privacy gate for every agent

- agent-guard exposes health, scan, and mask endpoints on loopback only.
- IDEs, MCP tools, and automations can check a resource immediately before attachment or transfer.
- The scan API returns categories, severity, and counts—never the raw findings.
- Passive file watching can create a local, value-free operational trail.

LIVE COMMAND
./demos/05-agent-
sidecar.sh

NETWORK BOUNDARY
IPv4 loopback only
Zero outbound calls

Why this matters for government

- Data sovereignty — sensitive content is inspected and transformed locally.
- Data minimisation — agents receive the minimum useful representation, not the full source record.
- Purpose limitation — Guardian incorporates the declared task into the release plan.
- Traceability — value-free reports and decision events support review without creating a second leak.
- Adoption — CLI, sidecar, and agent-rule integration fit existing developer and automation workflows.

PREVENTIVE CONTROL
Before transfer

DETECTIVE EVIDENCE
After decision

ARABIC READY

Trust boundaries and honest limitations

- Kakashi makes no outbound calls and submits no telemetry; database drivers connect only where the operator directs.
- Pattern detection is not infallible: organisations still need access control, classification, retention, and incident response.
- Audit modes reveal plaintext by design and must remain in a local human-only terminal.
- Kakashi supports PDPL-oriented controls; it does not certify compliance or replace legal review.
- An integration must honour Guardian decisions—the control is strongest when enforced in the transfer path.

SAFE CLAIM
Reduces accidental
disclosure risk.

UNSAFE CLAIM
“No leak is ever possible.”

Proof, not promises

- Six standalone demos can run in any order and rebuild only an ignored synthetic workspace.
- Each case asserts source-value removal, safe reports, stable database tokens, or fail-closed decisions.
- The complete project test suite covers detectors, formats, reporting, databases, sidecar, and Guardian.
- Kakashi scans the demo and presentation assets before they are uploaded to GitHub.

PRE-MEETING

```
npm test  
./demos/run-all.sh
```

LIVE

Run one numbered case
at a time

Practical government rollout

- PILOT — choose one low-risk workflow and measure findings, false positives, utility, and operator friction.
- POLICY — define destinations, approval owners, permitted classes, retention, and exception handling.
- INTEGRATE — place Guardian or the sidecar directly in the file/database-to-agent transfer path.
- OPERATE — review value-free evidence, tune policy, test upgrades, and train users on unsafe audit options.
- EXPAND — move from a team workflow to shared repositories, data platforms, and managed agent environments.

30 DAYS
Pilot + baseline

60 DAYS
Policy + integration

90 DAYS
Operate + expand

Tomorrow's live path

- 1 File protection
- 2 Folder compliance and Arabic report
- 3 Database result masking
- 4 Guardian transformation and human approval
- 5 Loopback agent sidecar
- 6 Operational evidence and Arabic CLI

START HERE
./demos/01-file-
protection.sh

RUNBOOK
demos/
PRESENTER_RUNBOOK
.md

GOAL
Safe AI utility without raw-
data exposure